

Computational Statistics

Hyperspherical VAE

Victor Deng Inès Vati

École Normale Supérieure Paris-Saclay, Master MVA

10th JAN 2024



Table of contents

1. Introduction
2. Sampling method
3. Reparameterization Trick
4. Experiments on link prediction
5. Conclusion and Discussion



Table of contents

1. Introduction
2. Sampling method
3. Reparameterization Trick
4. Experiments on link prediction
5. Conclusion and Discussion



2018 paper from Tim R. Davidson *et al.* [DFDC⁺]

- Replacing the Gaussian prior and approximate posterior distributions with a von Mises-Fisher distribution
- Goal: better model data with a hyperspherical latent structure
- Various experiments, where the \mathcal{S} -VAE (von Mises-Fisher distributions) often outperforms the \mathcal{N} -VAE (Gaussian distributions) in low dimensions



Table of contents


1. Introduction
2. Sampling method
3. Reparameterization Trick
4. Experiments on link prediction
5. Conclusion and Discussion



Démontrer que la méthode de sampling marche




Sampling w from $g(w|\kappa, \theta)$



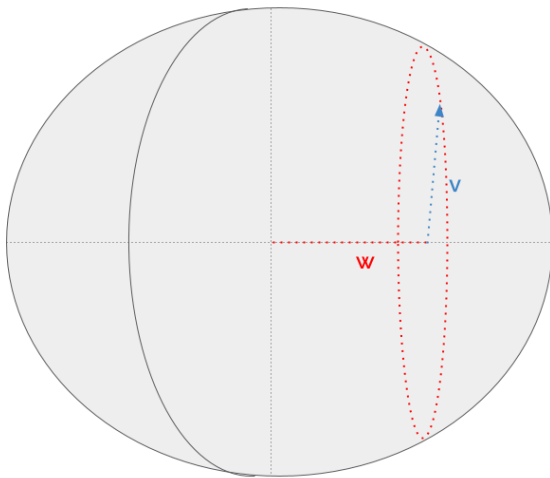
S^2 : unit sphere in \mathbb{R}^3



Sampling w



Sampling w



Sampling w




Table of contents

1. Introduction
2. Sampling method
3. Reparameterization Trick
4. Experiments on link prediction
5. Conclusion and Discussion



Reparameterization Trick

The authors use a reparameterization trick that has been extended to distributions that can be sampled using rejection sampling [NRLB20].

Algorithm 1 Reparameterized Rejection Sampling (from [NRLB20])

```
1:  $i \leftarrow 0$ 
2: repeat
3:    $i \leftarrow i + 1$ 
4:   Propose  $\varepsilon_i \sim s(\varepsilon)$ 
5:   Simulate  $u_i \sim \mathcal{U}[0, 1]$ 
6: until  $u_i < \frac{g(h(\varepsilon_i, \theta); \theta)}{r(h(\varepsilon_i, \theta); \theta)}$ 
7: return  $\varepsilon_i$ 
```



By noting $\pi(\varepsilon|\theta)$ the distribution of the resulting ε , we have

$$\nabla_{\theta} \mathbb{E}_{g(\varepsilon|\theta)}[\dots] = \mathbb{E}_{\pi(\varepsilon|\theta)}[\dots] = \mathbb{E}_{(\varepsilon_i, U_i)_i}[\dots]$$

Problem: $(\varepsilon_i, U_i)_{i \in \mathbb{N}}$ is not a random variable (it is a stochastic process)
No reference to a convergence proof in [DFDC⁺, NRLB20, PBJ12, MG14]



Table of contents

1. Introduction
2. Sampling method
3. Reparameterization Trick
4. Experiments on link prediction
5. Conclusion and Discussion



Experiments on link prediction

- Link prediction on a graph dataset: given a graph with some edges removed, predict the likelihood for each pair of nodes to be connected by an edge
- Cora dataset [MNRS00]: 2708 publications, 5429 links, 1433-dimensional feature vectors
- Using a Variational Graph Auto-Encoder [KW]: a variational encoder which uses a graph neural network (GNN) as encoder
- Reconstruction loss:

$$\mathbb{E}_{q(\mathbf{Z}|\mathbf{X},\mathbf{A})}(\log p(\mathbf{A}|\mathbf{Z})) \quad \text{where } p(\mathbf{A}|\mathbf{Z}) = \prod_{i=1}^N \prod_{j=1}^N p(A_{i,j}|\mathbf{z}_i, \mathbf{z}_j)$$

- Negative sampling: in the sum $\sum_{i,j} \log p(A_{i,j}|\mathbf{z}_i, \mathbf{z}_j)$, keep all positive edges and one randomly sampled negative edge per positive edge



reproduire l'experience

- data (Inès)
- ~~implémenter les modèles~~ (Victor VGAE)
- gradients pour le hyperspherical VAE (Inès)
- courbes d'entraînement dans le cas normal (Victor)
- entraînement et evaluation



Table of contents

1. Introduction
2. Sampling method
3. Reparameterization Trick
4. Experiments on link prediction
5. Conclusion and Discussion








Conclusion and Discussion

- Quite meaningful contribution in low dimensions
- Algorithm not really useful in high dimensions, due to vanishing surface problem and soap bubble effect of the \mathcal{N} -VAE
- Much less variance parameters (1 vs. d for \mathcal{N} -VAE), so possibly less expressivity
- vérifier différentes dimensions de l'espace latent
- et algo vraiment utile en petite ou moyenne dimension ?



References

-  Tim R. Davidson, Luca Falorsi, Nicola De Cao, Thomas Kipf, and Jakub M. Tomczak.
Hyperspherical variational auto-encoders.
-  Thomas N. Kipf and Max Welling.
Variational graph auto-encoders.
-  Andriy Mnih and Karol Gregor.
Neural variational inference and learning in belief networks, 2014.
-  Andrew Kachites McCallum, Kamal Nigam, Jason Rennie, and Kristie Seymore.
Automating the construction of internet portals with machine learning.
Information Retrieval, 3(2):127–163, 2000.
-  Christian A. Naesseth, Francisco J. R. Ruiz, Scott W. Linderman, and David M. Blei.
Reparameterization gradients through acceptance-rejection sampling algorithms, 2020.

